

Predictive Analysis Results and Code

Andrew Mashhadi

Since the R script became an out of control mess, I made this so we can see the code, results, and output pretty easily.

All Libraries Needed

```
library(rpart)
library(partykit)
library(tidyverse)
library(randomForest)
library(tidyverse)
library(tidymodels)
library(vip)
library(ggparty)
```

Predicting imdb-rating with single tree and rf

```
set.seed(777)

imdb_details_extd2 <- read.csv("C:\\Users\\amiro\\Desktop\\Statistics 405\\Week 5\\Final_Project_Brains")
imdb_details_extd2$star_power <- log(imdb_details_extd2$star_power+1)
imdb_details_extd2$wr_pop <- log(imdb_details_extd2$wr_pop+1)

## randomly select genres if more than one
tt<- lapply(imdb_details_extd2$genres, strsplit, ", ")
r_genre <- c()
for (i in 1:length(tt)) {


  if (identical(tt[[i]][[1]], character(0))) {
    name <- "None"
  } else {
    name <- sample(tt[[i]][[1]], 1)
  }
  r_genre <- c(r_genre, name)
}

imdb_details_extd2$genres <- r_genre

## simple single tree
tr <- rpart(imDbRating ~ runtime+genres+rating+dir_pop_fac+co_size+star_power+
            wr_pop+release_period+budget_adj,
```

```
imdb_details_extd2)

## print plot to help choose cp
plotcp(tr)
```



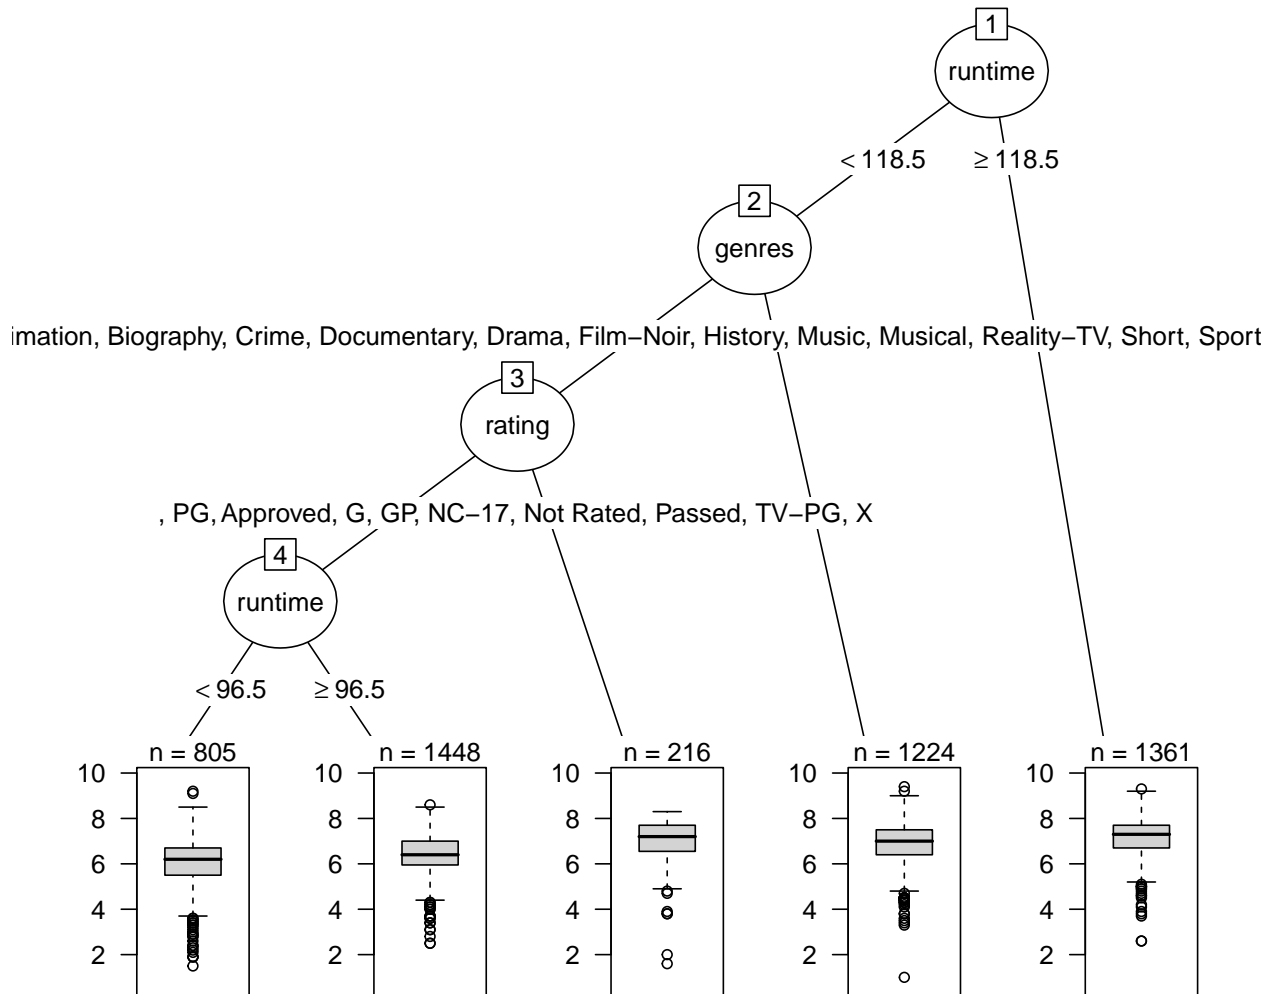
```
## prune the tree
tr_2 <- rpart::prune(tr, cp = 0.014)
print(tr_2)

## n= 5054
##
## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 5054 4737.3100 6.703720
##    2) runtime< 118.5 3693 3462.9790 6.530301
##      4) genres=Action,Adventure,Comedy,Family,Fantasy,Horror,Mystery,None,Romance,Sci-Fi,Thriller 24
##        8) rating=,PG,PG-13,R,TV-14,TV-MA,Unrated 2253 2013.4540 6.295473
##          16) runtime< 96.5 805 948.4268 6.047329 *
```

```
##      17) runtime>=96.5 1448  987.9022 6.433425 *
##      9) rating=Approved,G,GP,NC-17,Not Rated,Passed,TV-PG,X 216  216.8998 7.000926 *
##      5) genres=Animation,Biography,Crime,Documentary,Drama,Film-Noir,History,Music,Musical,Reality-TV,Short,Sport
##      3) runtime>=118.5 1361  861.8999 7.174284 *
```

```
## plot the pruned tree
```

```
plot(as.party(tr_2), tp_args = list(id = FALSE))
```



```
set.seed(777)
samp <- sample(5090)
rf_errors <- matrix(0, nrow=1, ncol=10)
lm_errors <- matrix(0, nrow=1, ncol=10)
pt_errors <- matrix(0, nrow=1, ncol=10)

for(k in 1:10){
  from <- 1 + (k-1)*509
  to <- 509*k # we will lose the last 8 observations
  test <- na.omit(imdb_details_extd2[samp[from:to],])
  train <- imdb_details_extd2[samp[-(from:to)],]
```

```

## rf
imdb_rf <- randomForest(imDbRating ~ runtime+genres+rating+dir_pop_fac+
                        co_size+star_power+wr_pop+release_period+budget_adj,
                        data = imdb_details_extd2,
                        mtry = 3,
                        na.action = na.omit)

# lm
imdb_lm <- lm(imDbRating ~runtime+rating+dir_pop_fac+
              co_size+star_power+wr_pop+budget_adj,
              data = imdb_details_extd2,
              na.action = na.omit)

## single pruned tree
tr <- rpart(imDbRating ~runtime+genres+rating+dir_pop_fac+
            co_size+star_power+wr_pop+release_period+budget_adj,
            data = imdb_details_extd2)
pt <- rpart::prune(tr, cp = 0.014)

## calc errors for this fold
rf_errors[k] <- mean((test$imDbRating - predict(imdb_rf, test))^2)
lm_errors[k] <- mean((test$imDbRating - predict(imdb_lm, test))^2)
pt_errors[k] <- mean((test$imDbRating - predict(pt, test))^2)
}

## compare K-fold MSEs
mean(rf_errors)

## [1] 0.68611
mean(lm_errors)

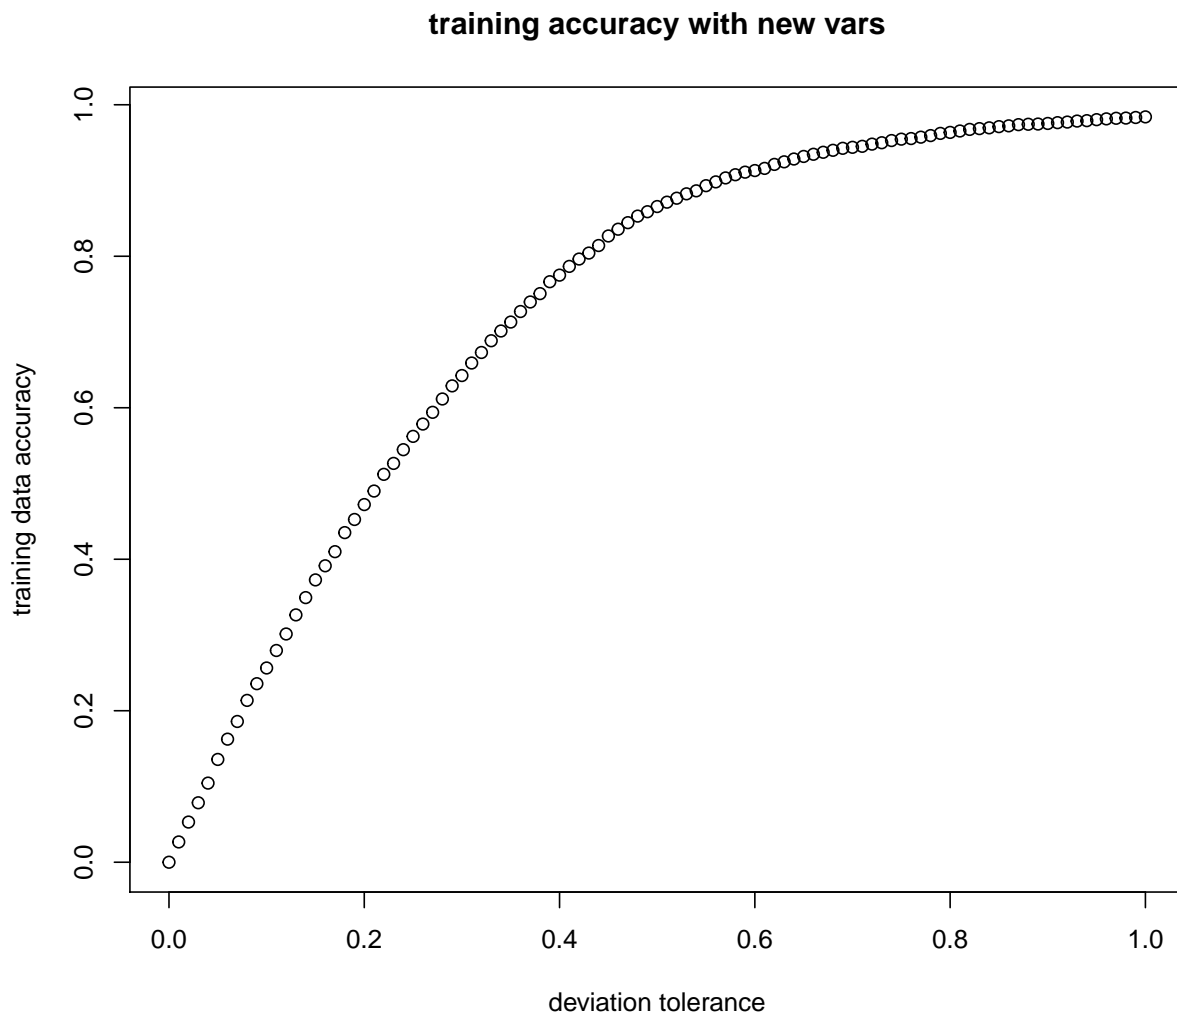
## [1] 0.6849048
mean(pt_errors)

## [1] 0.7465697
set.seed(777)
## obtain prediction accuracy
imdb_details_no_NA <- na.omit(imdb_details_extd2)
imdb_rf = randomForest(imDbRating ~runtime+genres+rating+dir_pop_fac+
                        co_size+star_power+wr_pop+release_period+budget_adj,
                        data = imdb_details_no_NA,
                        mtry = 3)

## compare the predictions to the data
tr_comp <- data.frame(imDbRating=imdb_details_no_NA$imDbRating, predictedRating=predict(imdb_rf, imdb_d

## approximate training accuracy
devs <- abs(tr_comp$imDbRating - tr_comp$predictedRating)
close_enoughs <- function(x) sum(devs <= x) / length(devs)
x <- seq(from=0,to=1,by=0.01)
plot(x, sapply(x, close_enoughs), main="training accuracy with new vars", xlab="deviation tolerance", y

```

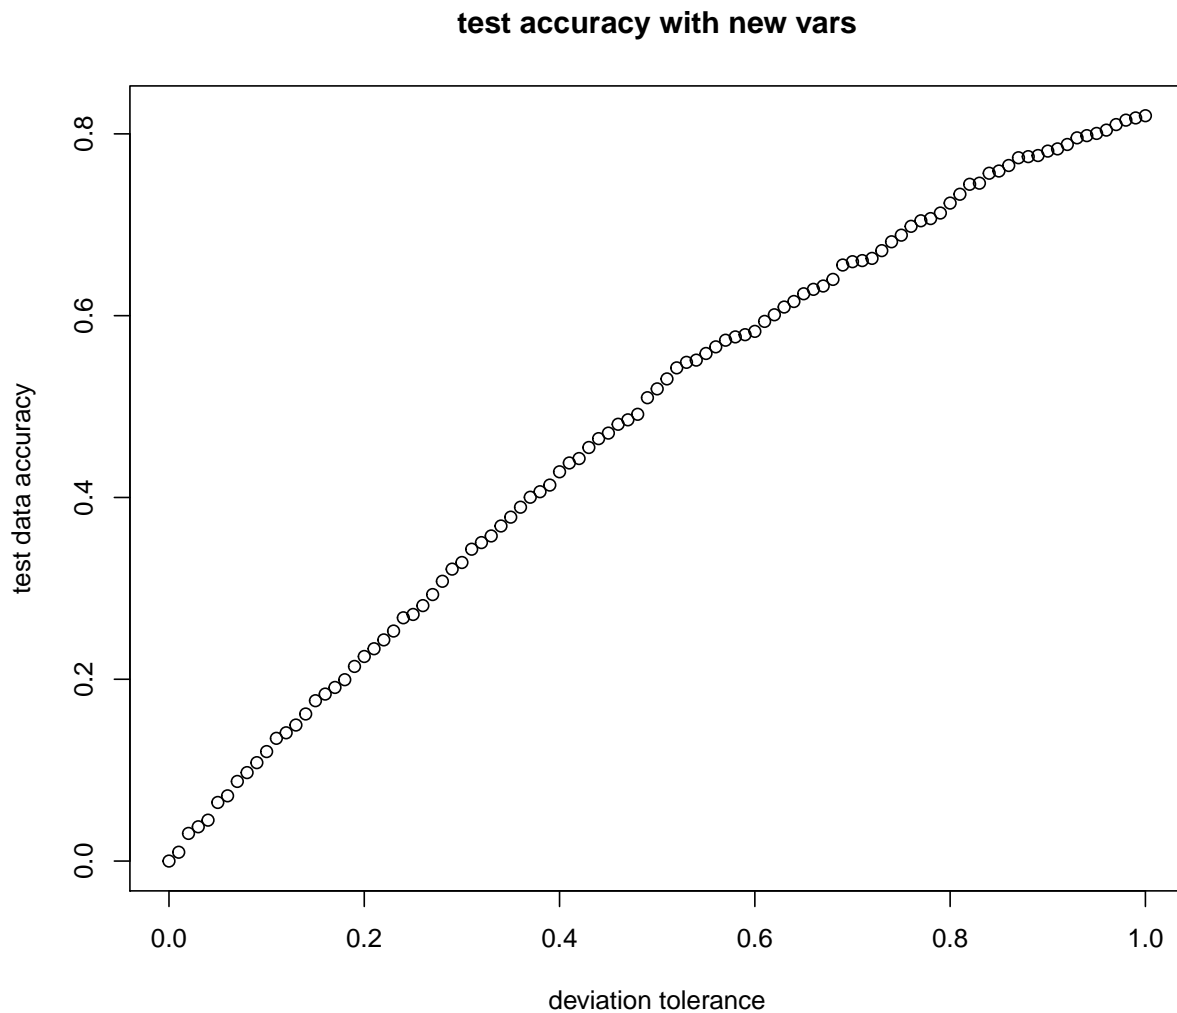


```
## approximate test accuracy
samp <- sample(5098)
train <- na.omit(imdb_details_extd2[samp[1:4000],])
test  <- na.omit(imdb_details_extd2[samp[4001:5098],])

imdb_rf = randomForest(imDbRating ~runtime+genres+rating+dir_pop_fac+
                        co_size+star_power+wr_pop+release_period+budget_adj,
                        data = train,
                        mtry = 3)

## compare the predictions to the data
comp <- data.frame(imDbRating=test$imDbRating, predictedRating=predict(imdb_rf, test))

devs <- abs(comp$imDbRating - comp$predictedRating)
close_enoughs <- function(x) sum(devs <= x)/ length(devs)
x <- seq(from=0,to=1,by=0.01)
plot(x, sapply(x, close_enoughs), main="test accuracy with new vars", xlab="deviation tolerance", ylab=
```



Predicting imdb-rating with TidyModels

```
set.seed(777)
trees_df <- filter(imdb_details_extd2, type == "Movie") %>%
  na.omit()

trees_split <- initial_split(trees_df)
trees_train <- training(trees_split)
trees_test <- testing(trees_split)

# build recipe (just instructions)
tree_rec <- recipe(imDbRating ~ runtime+genres+rating+dir_pop_fac+
  co_size+star_power+wr_pop+release_period+budget_adj,
  data = trees_train) %>%
  step_other(genres, threshold = 0.03) %>%
```

```

    step_unknown(genres) %>%
    step_other(rating, threshold = 0.05) %>%
    step_unknown(rating) %>%
    step_dummy(all_nominal(), -all_outcomes())

# prep actually uses the data
tree_prep <- prep(tree_rec)
juiced <- juice(tree_prep)

# run the below to check the step_other results (doesn't work if step_dummy already used)
# juiced %>% count(genres, sort = T)

# report details
summary(tree_rec)

## # A tibble: 10 x 4
##   variable      type    role    source
##   <chr>         <chr>  <chr>   <chr>
## 1 runtime      numeric predictor original
## 2 genres       nominal predictor original
## 3 rating       nominal predictor original
## 4 dir_pop_fac  numeric predictor original
## 5 co_size      numeric predictor original
## 6 star_power   numeric predictor original
## 7 wr_pop       numeric predictor original
## 8 release_period nominal predictor original
## 9 budget_adj   numeric predictor original
## 10 imdbRating  numeric outcome  original

# build model for tuning
tune_spec <- rand_forest(
  mtry = tune(),
  trees = tune()
) %>%
  set_mode("regression") %>%
  set_engine("ranger")

tune_wf <- workflow() %>%
  add_recipe(tree_rec) %>%
  add_model(tune_spec)

# create a set of cross-validation resamples to use for tuning
trees_folds <- vfold_cv(trees_train)


# choose 10 grid points automatically
tune_res <- tune_grid(
  tune_wf,
  resamples = trees_folds,
  grid = 10
)

## i Creating pre-processing data to finalize unknown parameter: mtry
## Warning: package 'ranger' was built under R version 4.0.5

```

```
### rmse plot for tuning mtry and number of trees
```

```
tune_res %>%
  collect_metrics() %>%
  filter(.metric == "rmse") %>%
  select(mean, trees, mtry) %>%
  pivot_longer(trees:mtry,
               values_to = "value",
               names_to = "parameter")
) %>%
ggplot(aes(value, mean, color = parameter)) +
  geom_point(show.legend = FALSE) +
  facet_wrap(~parameter, scales = "free_x") +
  labs(x = NULL, y = "rmse")
```



```
### rsq plot for tuning mtry and number of trees
```


```
tune_res %>%
  collect_metrics() %>%
  filter(.metric == "rsq") %>%
  select(mean, trees, mtry) %>%
```



```

pivot_longer(trees:mtry,
              values_to = "value",
              names_to = "parameter")
) %>%
ggplot(aes(value, mean, color = parameter)) +
geom_point(show.legend = FALSE) +
facet_wrap(~parameter, scales = "free_x") +
labs(x = NULL, y = "Rsquared")

```



```
## looks like 3 for mtry and 1000-1500 for trees could work best
```

```
### taking a closer look now
```


```

rf_grid <- grid_regular(
  mtry(range = c(2, 5)),
  trees(range = c(1000, 1500)),
  levels = 4
)

```

```
regular_res <- tune_grid(
  tune_wf,
  resamples = trees_folds,
  grid = rf_grid
)

## rmse plot for tuning mtry and number of trees
regular_res %>%
  collect_metrics() %>%
  filter(.metric == "rmse") %>%
  mutate(trees = factor(trees)) %>%
  ggplot(aes(mtry, mean, color = trees)) +
  geom_line(alpha = 0.5, size = 1.5) +
  geom_point() +
  labs(y = "rmse")
```




```
## rsq plot for tuning mtry and number of trees
regular_res %>%
  collect_metrics() %>%
```

```

filter(.metric == "rsq") %>%
mutate(trees = factor(trees)) %>%
ggplot(aes(mtry, mean, color = trees)) +
geom_line(alpha = 0.5, size = 1.5) +
geom_point() +
labs(y = "rsq")

```



```
## looks like 3 for mtry and 1166 for trees is optimal
```

```
## build model with tuned params
```

```

final_rf <- rand_forest(
  mtry = 3,
  trees = 1166,
) %>%
  set_mode("regression") %>%
  set_engine("ranger")

```


```
# checking out importance plots
```

```
final_rf %>%
```

```

set_engine("ranger", importance = "permutation") %>%
fit(imDbRating ~ .,
    data = juice(tree_prep)
) %>%
vip(geom = "point")

```



```
### view metrics
```

```

final_wf <- workflow() %>%
  add_recipe(tree_rec) %>%
  add_model(final_rf)

```

```

final_res <- final_wf %>%
  last_fit(trees_split)

```

```

final_res %>%
  collect_metrics()

```

```
## # A tibble: 2 x 4
```

```
## .metric .estimator .estimate .config
## <chr> <chr> <dbl> <chr>
## 1 rmse standard 0.816 Preprocessor1_Model11
## 2 rsq standard 0.326 Preprocessor1_Model11
```

Predicting Oscar-Nomination with single tree and rf

```
set.seed(777)


imdb_details_extd2 <- read.csv("C:\\Users\\amiro\\Desktop\\Statistics 405\\Week 5\\Final_Project_Brains")
imdb_details_extd2$star_power <- log(imdb_details_extd2$star_power+1)
imdb_details_extd2$wr_pop <- log(imdb_details_extd2$wr_pop+1)

## randomly select genres if more than one
tt<- lapply(imdb_details_extd2$genres, strsplit, ", ")
r_genre <- c()
for (i in 1:length(tt)) {

  if (identical(tt[[i]][[1]], character(0))) {
    name <- "None"
  } else {
    name <- sample(tt[[i]][[1]], 1)
  }
  r_genre <- c(r_genre, name)
}

imdb_details_extd2$genres <- r_genre
imdb_details_extd2$oscar_nom <- as.factor(imdb_details_extd2$oscar_nom)

## simple single tree
tr <- rpart(oscar_nom ~ runtime+genres+rating+dir_pop_fac+co_size+star_power+
            wr_pop+release_period+budget_adj, data=imdb_details_extd2)
## print plot to help choose cp
plotcp(tr)
```



```
## prune the tree
```

```
tr_2 <- rpart::prune(tr, cp = 0.013)
```

```
print(tr_2)
```

```
## n= 5054
```

```
##
```

```
## node), split, n, loss, yval, (yprob)
```

```
##      * denotes terminal node
```

```
##
```

```
## 1) root 5054 1219 0 (0.75880491 0.24119509)
```

```
## 2) runtime< 117.5 3593 585 0 (0.83718341 0.16281659) *
```

```
## 3) runtime>=117.5 1461 634 0 (0.56605065 0.43394935)
```

```
## 6) rating=M,NC-17,Not Rated,TV-14,TV-G,TV-MA,TV-PG,Unrated 152 10 0 (0.93421053 0.06578947) *
```

```
## 7) rating=Approved,G,GP,M/PG,Passed,PG,PG-13,R,X 1309 624 0 (0.52330023 0.47669977)
```

```
## 14) runtime< 147.5 1110 478 0 (0.56936937 0.43063063)
```

```
## 28) release_period=aug_to_oct,feb_to_apr,may_to_jul 750 272 0 (0.63733333 0.36266667)
```

```
## 56) rating=GP,PG-13,R 633 201 0 (0.68246445 0.31753555) *
```

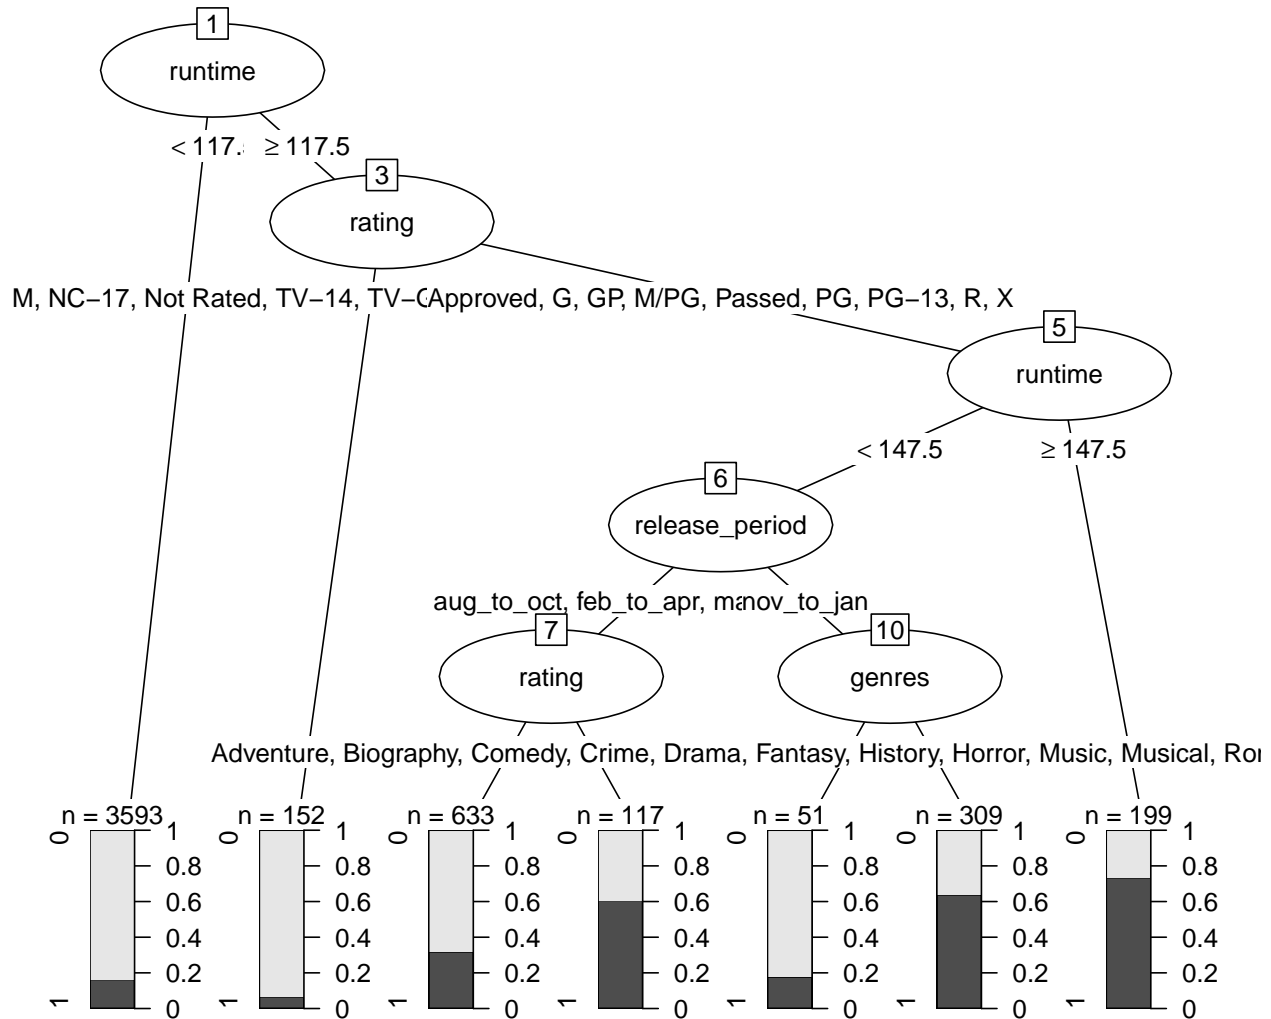
```
## 57) rating=Approved,G,Passed,PG,X 117 46 1 (0.39316239 0.60683761) *
```

```
## 29) release_period=nov_to_jan 360 154 1 (0.42777778 0.57222222)
```

```
##          58) genres=Action,Animation,Family,Mystery,Western 51    9 0 (0.82352941 0.17647059) *
##          59) genres=Adventure,Biography,Comedy,Crime,Drama,Fantasy,History,Horror,Music,Musical,Romance 10 0 (0.82352941 0.17647059) *
##          15) runtime>=147.5 199    53 1 (0.26633166 0.73366834) *
```

```
## plot the pruned tree
```

```
plot(as.party(tr_2), tp_args = list(id = FALSE))
```



```
##### COMPARING K-FOLD accuracy FOR MULTIPLE MODELS #####
```

```
samp <- sample(5090)
rf_errors <- matrix(0, nrow=1, ncol=10)
pt_errors <- matrix(0, nrow=1, ncol=10)

for(k in 1:10){
  from <- 1 + (k-1)*509
  to <- 509*k # we will lose the last 8 observations
  test <- na.omit(imdb_details_extd2[samp[from:to],])
  train <- imdb_details_extd2[samp[-(from:to)],]
```

```
## rf
```

```

imdb_rf <- randomForest(oscar_nom ~ runtime+genres+rating+dir_pop_fac+co_size+star_power+
                        wr_pop+release_period+budget_adj,
                        data = imdb_details_extd2,
                        mtry = 3,
                        na.action = na.omit)

## single pruned tree
tr <- rpart(oscar_nom ~ runtime+genres+year+rating+dir_pop_fac+co_size+star_power+
            wr_pop+release_period+budget_adj,
            data = imdb_details_extd2)
pt <- rpart::prune(tr, cp = 0.011)

## calc pred accuracy for this fold
rf_errors[k] <- mean(test$oscar_nom==predict(imdb_rf, test))
pt_preds <- apply(predict(pt, test), 1, which.max) - 1
pt_errors[k] <- mean(test$oscar_nom==pt_preds)
}

## compare K-fold prediction accuracy
mean(rf_errors)

## [1] 0.9170455
mean(pt_errors)

## [1] 0.7890319

```

Predicting Oscar Nominations with TidyModels

```

set.seed(777)

trees_df <- filter(imdb_details_extd2, type == "Movie") %>%
  na.omit()

trees_split <- initial_split(trees_df)
trees_train <- training(trees_split)
trees_test <- testing(trees_split)

# build recipe (just instructions)

tree_rec <- recipe(oscar_nom ~ runtime+genres+rating+dir_pop_fac+co_size+star_power+
                   wr_pop+release_period+budget_adj, data = trees_train) %>%
  step_other(genres, threshold = 0.03) %>%
  step_unknown(genres) %>%
  step_other(rating, threshold = 0.05) %>%
  step_unknown(rating) %>%
  step_dummy(all_nominal(), -all_outcomes())

# prep actually uses the data

```



```

tree_prep <- prep(tree_rec)
juiced <- juice(tree_prep)

# check the step_other results
# juiced %>% count(genres, sort = T)

summary(tree_rec)

## # A tibble: 10 x 4
##   variable      type    role    source
##   <chr>         <chr>  <chr>   <chr>
## 1 runtime      numeric predictor original
## 2 genres       nominal predictor original
## 3 rating       nominal predictor original
## 4 dir_pop_fac  numeric predictor original
## 5 co_size      numeric predictor original
## 6 star_power   numeric predictor original
## 7 wr_pop       numeric predictor original
## 8 release_period nominal predictor original
## 9 budget_adj   numeric predictor original
## 10 oscar_nom   nominal outcome  original

# build model
tune_spec <- rand_forest(
  mtry = tune(),
  trees = tune()
) %>%
  set_mode("classification") %>%
  set_engine("ranger")

tune_wf <- workflow() %>%
  add_recipe(tree_rec) %>%
  add_model(tune_spec)

# create a set of cross-validation resamples to use for tuning
trees_folds <- vfold_cv(trees_train)

#doParallel::registerDoParallel() #try me

# choose 10 grid points automatically
tune_res <- tune_grid(
  tune_wf,
  resamples = trees_folds,
  grid = 10
)


## i Creating pre-processing data to finalize unknown parameter: mtry
### roc_auc plot for tuning mtry and number of trees
tune_res %>%
  collect_metrics() %>%
  filter(.metric == "roc_auc") %>%
  dplyr::select(mean, trees, mtry) %>%
  pivot_longer(trees:mtry,
    values_to = "value",

```

```

names_to = "parameter"
) %>%
ggplot(aes(value, mean, color = parameter)) +
geom_point(show.legend = FALSE) +
facet_wrap(~parameter, scales = "free_x") +
labs(x = NULL, y = "AUC")

```



```

### accuracy plot for tuning mtry and number of trees


```

```

tune_res %>%
  collect_metrics() %>%
  filter(.metric == "accuracy") %>%
  dplyr::select(mean, trees, mtry) %>%
  pivot_longer(trees:mtry,
               values_to = "value",
               names_to = "parameter"
  ) %>%
  ggplot(aes(value, mean, color = parameter)) +
  geom_point(show.legend = FALSE) +
  facet_wrap(~parameter, scales = "free_x") +

```

```
labs(x = NULL, y = "Accuracy")
```



```
## looks like 3 for mtry and 1000-1500 for trees could work best
```


```
### taking a closer look now
```

```
rf_grid <- grid_regular(  
  mtry(range = c(3, 10)),  
  trees(range = c(1000, 1500)),  
  levels = 6  
)
```

```
regular_res <- tune_grid(  
  tune_wf,  
  resamples = trees_folds,  
  grid = rf_grid  
)
```


```
## AUC plot for tuning mtry and number of trees
```

```
regular_res %>%
  collect_metrics() %>%
  filter(.metric == "roc_auc") %>%
  mutate(trees = factor(trees)) %>%
  ggplot(aes(mtry, mean, color = trees)) +
  geom_line(alpha = 0.5, size = 1.5) +
  geom_point() +
  labs(y = "AUC")
```



accuracy plot for tuning mtry and number of trees


```
regular_res %>%
  collect_metrics() %>%
  filter(.metric == "accuracy") %>%
  mutate(trees = factor(trees)) %>%
  ggplot(aes(mtry, mean, color = trees)) +
  geom_line(alpha = 0.5, size = 1.5) +
  geom_point() +
  labs(y = "Accuracy")
```



```
## looks like 7 for mtry and 1250 for trees is optimal

## build model with tuned params
final_rf <- rand_forest(
  mtry = 7,
  trees = 1250,
) %>%
  set_mode("classification") %>%
  set_engine("ranger")

# checking out importance plots (measure of how much worse predictions get w/o the var)
final_rf %>%
  set_engine("ranger", importance = "permutation") %>%
  fit(oscar_nom ~ .,
      data = juice(tree_prep)
  ) %>%
  vip(geom = "point")
```



```
### view metrics
```

```
final_wf <- workflow() %>%
  add_recipe(tree_rec) %>%
  add_model(final_rf)
```

```
final_res <- final_wf %>%
  last_fit(trees_split)
```

```
final_res %>%
  collect_metrics()
```

```
## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>      <dbl> <chr>
## 1 accuracy binary      0.801 Preprocessor1_Model11
## 2 roc_auc  binary      0.800 Preprocessor1_Model11
```

Precision-Response AUC and Confusion Matrix for Oscar Nominations (with TidyModels)

This should be compared with Dan's Logistic Regression model.

```
set.seed(777)

## check pr_auc
## good ref: https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-imbalanced-c
## want to keep precision AND recall high as threshold increases

fit_rf_oscar_nom <-
  final_wf %>%
  fit(trees_train)

pred_rf_oscar_nom <-
  predict(fit_rf_oscar_nom, trees_test) %>%
  bind_cols(predict(fit_rf_oscar_nom, trees_test, type = "prob")) %>%
  bind_cols(trees_test %>% dplyr::select(oscar_nom))

pred_rf_oscar_nom %>%
  pr_auc(event_level = "second", truth = oscar_nom, .pred_1)

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 pr_auc binary      0.578

###
pred_v_real <- final_res %>%
  dplyr::select(.predictions) %>%
  unnest(.predictions) %>%
  as_tibble() %>%
  dplyr::select(.pred_class, oscar_nom)

conf_mat(pred_v_real, oscar_nom, .pred_class)

##           Truth
## Prediction    0    1
##           0 662 142
##           1  48 101
```

Predicting Box Office Profits with Single Tree and RF

```
set.seed(777)

imdb_details_extd3 <- read.csv("C:\\Users\\amiro\\Desktop\\Statistics 405\\Week 5\\Final_Project_Brains")
imdb_details_extd3$star_power <- log(imdb_details_extd3$star_power+1)
imdb_details_extd3$wr_pop <- log(imdb_details_extd3$wr_pop+1)
imdb_details_extd3$inf_adjusted_gp = imdb_details_extd3$inf_adjusted_gp^.25

## randomly select genres if more than one
```

```


tt<- lapply(imdb_details_extd3$genres, strsplit, ", ")
r_genre <- c()
for (i in 1:length(tt)) {

  if (identical(tt[[i]][[1]], character(0))) {
    name <- "None"
  } else {
    name <- sample(tt[[i]][[1]], 1)
  }
  r_genre <- c(r_genre, name)
}

imdb_details_extd3$genres <- r_genre

## simple single tree
tr <- rpart(inf_adjusted_gp ~ genres+rating+dir_pop_fac+budget_adj+runtime+release_period+
            co_size+star_power+wr_pop, data=imdb_details_extd3)
## print plot to help choose cp
plotcp(tr)

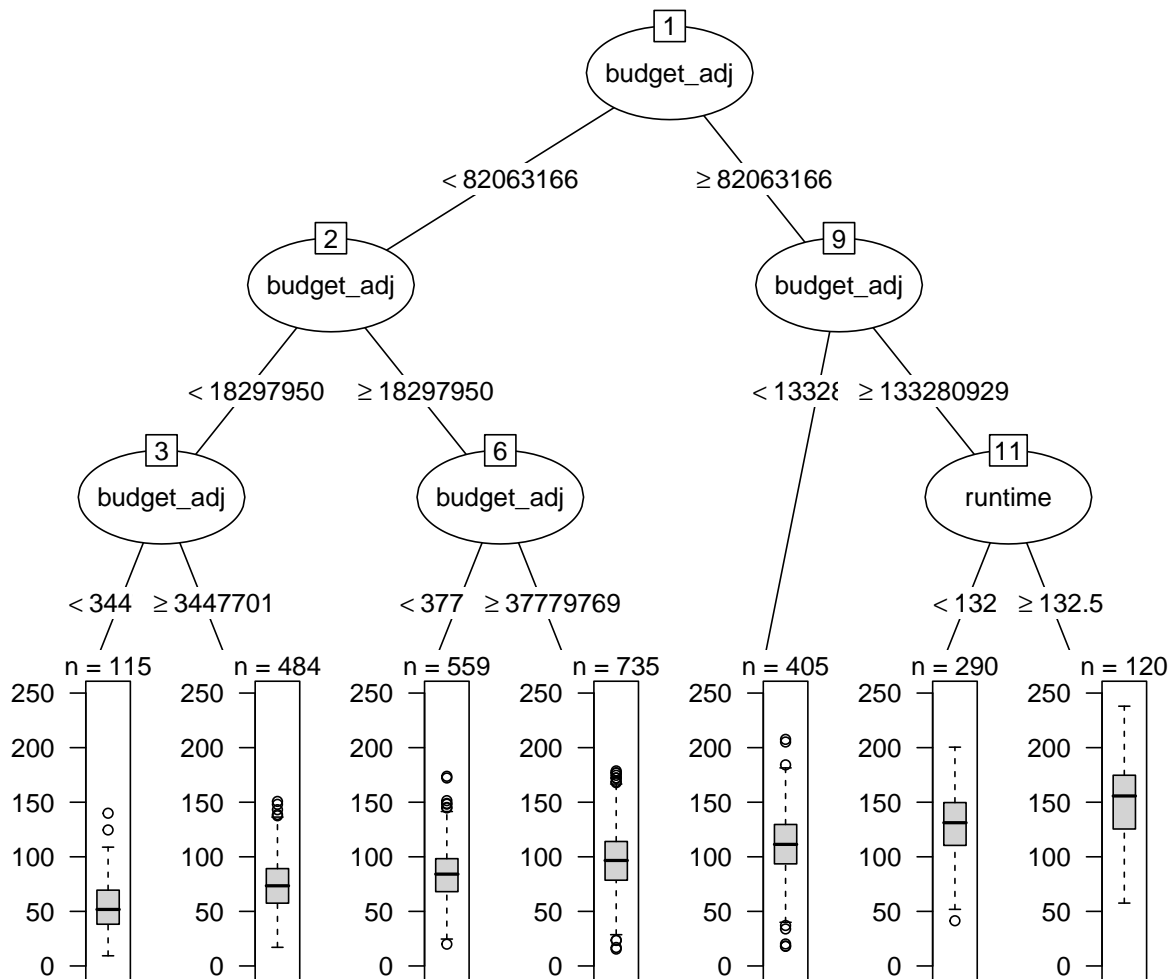
```

```
## prune the tree
tr_2 <- rpart::prune(tr, cp=0.012)
print(tr_2)

## n=2708 (700 observations deleted due to missingness)
##
## node), split, n, deviance, yval
##      * denotes terminal node
##
## 1) root 2708 3224647.0  96.51553
##    2) budget_adj< 8.206317e+07 1893 1498279.0  85.07092
##      4) budget_adj< 1.829795e+07 599  387245.1  70.44133 *
##      5) budget_adj>=1.829795e+07 1294  923488.2  91.84305
##        10) budget_adj< 3.777977e+07 559  326229.4  84.85289 *
##        11) budget_adj>=3.777977e+07 735  549171.3  97.15937 *
##    3) budget_adj>=8.206317e+07 815  902527.2 123.09790
##      6) budget_adj< 1.332809e+08 405  352237.6 110.89160 *
##      7) budget_adj>=1.332809e+08 410  430339.7 135.15540
##    14) runtime< 132.5 290  249979.6 128.76830 *
```

```
##      15) runtime>=132.5 120 139939.5 150.59080 *
## plot the pruned tree
plot(as.party(tr), tp_args = list(id = FALSE))
```



Predicting Box Office Profit with TidyModels

```
df <- filter(imdb_details_extd3, type == "Movie") %>%
  na.omit()

data_split <- initial_split(df)
data_train <- training(data_split)
data_test <- testing(data_split)

# build recipe (just instructions)
model_rec <- recipe(inf_adjusted_gp ~ budget_adj+runtime+genres+rating+dir_pop_fac+release_period+
```

```

        runtime+co_size+star_power+wr_pop,
        data = data_train) %>%
step_other(genres, threshold = 0.03) %>%
step_unknown(genres) %>%
step_other(rating, threshold = 0.05) %>%
step_unknown(rating) %>%
step_dummy(all_nominal(), -all_outcomes())

# prep actually uses the data
model_prep <- prep(model_rec)
juiced <- juice(model_prep)

# run the below to check the step_other results (doesnt work if step_dummy already used)
# juiced %>% count(genres, sort = T)

# report details
summary(model_rec)

## # A tibble: 10 x 4
##   variable      type    role    source
##   <chr>         <chr>  <chr>   <chr>
## 1 budget_adj    numeric predictor original
## 2 runtime       numeric predictor original
## 3 genres        nominal predictor original
## 4 rating        nominal predictor original
## 5 dir_pop_fac   numeric predictor original
## 6 release_period nominal predictor original
## 7 co_size       numeric predictor original
## 8 star_power    numeric predictor original
## 9 wr_pop        numeric predictor original
## 10 inf_adjusted_gp numeric outcome  original

# build model for tuning
rf_bo_profit_spec <- rand_forest(
  mtry = tune(),
  trees = tune()
) %>%
  set_mode("regression") %>%
  set_engine("ranger")

rf_tune_wf <- workflow() %>%
  add_recipe(model_rec) %>%
  add_model(rf_bo_profit_spec)

# create a set of cross-validation resamples to use for tuning
data_folds <- vfold_cv(data_train)

# choose 10 grid points automatically
rf_tune_res <- tune_grid(
  rf_tune_wf,
  resamples = data_folds,
  grid = 10
)


## i Creating pre-processing data to finalize unknown parameter: mtry

```

```
## ! Fold04: preprocessor 1/1, model 9/10: 25 columns were requested but there were 2...
## ! Fold07: preprocessor 1/1, model 9/10: 25 columns were requested but there were 2...
```

```
### rmse plot for tuning mtry and number of trees
```

```
rf_tune_res %>%
  collect_metrics() %>%
  filter(.metric == "rmse") %>%
  dplyr::select(mean, trees, mtry) %>%
  pivot_longer(trees:mtry,
               values_to = "value",
               names_to = "parameter")
) %>%
ggplot(aes(value, mean, color = parameter)) +
  geom_point(show.legend = FALSE) +
  facet_wrap(~parameter, scales = "free_x") +
  labs(x = NULL, y = "rmse")
```




```
### rsq plot for tuning mtry and number of trees
```

```
rf_tune_res %>%
```

```

collect_metrics() %>%
  filter(.metric == "rsq") %>%
  dplyr::select(mean, trees, mtry) %>%
  pivot_longer(trees:mtry,
               values_to = "value",
               names_to = "parameter")
) %>%
  ggplot(aes(value, mean, color = parameter)) +
  geom_point(show.legend = FALSE) +
  facet_wrap(~parameter, scales = "free_x") +
  labs(x = NULL, y = "Rsquared")

```



```
## looks like 3 for mtry and 1000-1500 for trees could work best
```

```
### taking a closer look now
```

```

rf_grid <- grid_regular(
  mtry(range = c(4, 10)),
  trees(range = c(800, 1200)),

```


```

    levels = 5
  )

rf_res <- tune_grid(
  rf_tune_wf,
  resamples = data_folds,
  grid = rf_grid
)


## rmse plot for tuning mtry and number of trees
rf_res %>%
  collect_metrics() %>%
  filter(.metric == "rmse") %>%
  mutate(trees = factor(trees)) %>%
  ggplot(aes(mtry, mean, color = trees)) +
  geom_line(alpha = 0.5, size = 1.5) +
  geom_point() +
  labs(y = "rmse")

```



```
## rsq plot for tuning mtry and number of trees
```

```
rf_res %>%
  collect_metrics() %>%
  filter(.metric == "rsq") %>%
  mutate(trees = factor(trees)) %>%
  ggplot(aes(mtry, mean, color = trees)) +
  geom_line(alpha = 0.5, size = 1.5) +
  geom_point() +
  labs(y = "rsq")
```



```
## looks like 7 for mtry and 1100 for trees is optimal
```


```
## build rf model with tuned params
```

```
final_rf <- rand_forest(
  mtry = 7,
  trees = 1100,
) %>%
  set_mode("regression") %>%
  set_engine("ranger")
```

```

# checking out importance plots
final_rf %>%
  set_engine("ranger", importance = "permutation") %>%
  fit(inf_adjusted_gp ~ .,
      data = juice(model_prep)
  ) %>%
  vip(geom = "point")

```



```

### view metrics

final_wf <- workflow() %>%
  add_recipe(model_rec) %>%
  add_model(final_rf)

final_res <- final_wf %>%
  last_fit(data_split)

final_res %>%

```



```
collect_metrics()
```

```
## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard      25.6  Preprocessor1_Model1
## 2 rsq     standard       0.457 Preprocessor1_Model1
```


Plot Accuracy For Box Office Profit Random Forest (transformed data back)

Results looks great ... when the data is transformed. After transforming back they look kind of like this.

```
fit_rf_oscar_nom <-
  final_wf %>%
  fit(data_train)

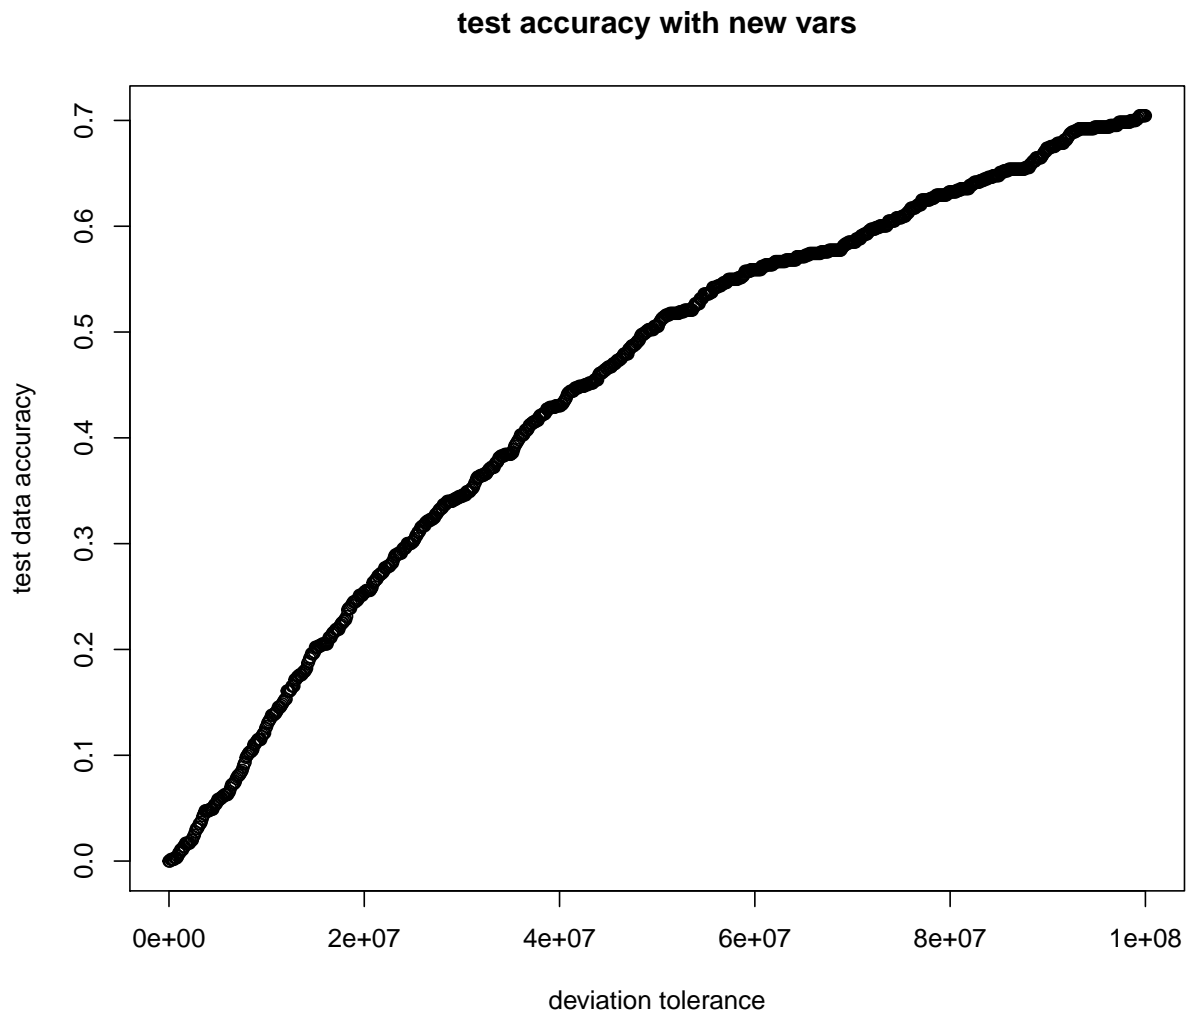
## compare the predictions to the data
tr_comp <- data.frame(true_gp=data_train$inf_adjusted_gp, predict(fit_rf_oscar_nom, data_train))

## approximate training accuracy
devs <- abs((tr_comp$true_gp)^4 - (tr_comp$.pred)^4)
close_enoughs <- function(x) sum(devs <= x) / length(devs)
x <- seq(from=0,to=100000000,by=100000)
plot(x, sapply(x, close_enoughs), main="training accuracy with new vars", xlab="deviation tolerance", ylab="accuracy")
```



```
## compare the predictions to the data
tr_comp <- data.frame(true_gp=data_test$inf_adjusted_gp, predict(fit_rf_oscar_nom, data_test))

## approximate test accuracy
devs <- abs(tr_comp$true_gp^4 - tr_comp$.pred^4)
x <- seq(from=0,to=100000000,by=100000)
plot(x, sapply(x, close_enoughs), main="test accuracy with new vars", xlab="deviation tolerance", ylab=
```



Compare Box Office Profit Random Forest with Linear Model

```
# build recipe (just instructions)
lm_model_rec <- recipe(inf_adjusted_gp ~ budget_adj+runtime+dir_pop_fac+release_period+
  runtime+co_size+star_power+wr_pop,
  data = data_train)

## build lm model with tuned params
final_lm <- linear_reg() %>%
  set_engine("lm") %>%
  set_mode("regression")

lm_final_wf <- workflow() %>%
  add_recipe(lm_model_rec) %>%
  add_model(final_lm)
```

```
lm_final_res <- lm_final_wf %>%
  last_fit(data_split)

lm_final_res %>%
  collect_metrics()

## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard      26.9   Preprocessor1_Model1
## 2 rsq     standard       0.402 Preprocessor1_Model1
```

Predicting GPM with Single Tree and RF

Pretty boring looking tree.

```
set.seed(777)

imdb_details_extd3 <- read.csv("C:\\Users\\amiro\\Desktop\\Statistics 405\\Week 5\\Final_Project_Brains\\imdb_details_extd3.csv")
imdb_details_extd3 <- filter(imdb_details_extd3, type == "Movie") %>%
  na.omit()
imdb_details_extd3$star_power <- log(imdb_details_extd3$star_power+1)
imdb_details_extd3$wr_pop <- log(imdb_details_extd3$wr_pop+1)
C <- min(imdb_details_extd3$gpm)
imdb_details_extd3$gpm = log(imdb_details_extd3$gpm+1-C)

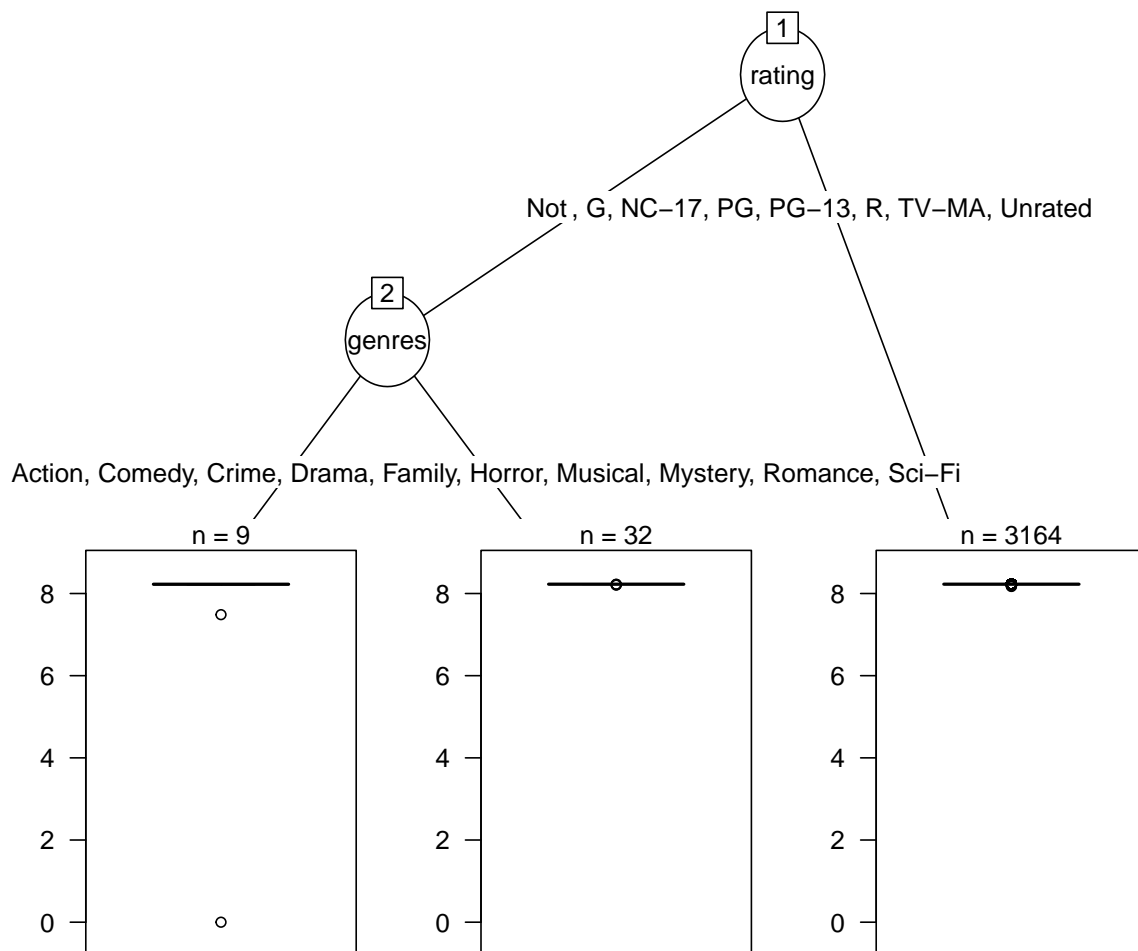
## randomly select genres if more than one
tt<- lapply(imdb_details_extd3$genres, strsplit, ", ")
r_genre <- c()
for (i in 1:length(tt)) {

  if (identical(tt[[i]][[1]], character(0))) {
    name <- "None"
  } else {
    name <- sample(tt[[i]][[1]], 1)
  }
  r_genre <- c(r_genre, name)
}

imdb_details_extd3$genres <- r_genre

## simple single tree
tr <- rpart(gpm ~ genres+rating+dir_pop_fac+budget_adj+runtime+release_period+
  co_size+star_power+wr_pop, data=imdb_details_extd3)

## plot the pruned tree
plot(as.party(tr), tp_args = list(id = FALSE))
```



Predicting GPM with Random Forest in TidyModels

Notice the RMSE is not that great given the transformation, and the R^2 is absolute trash. Not sure if I made a mistake or something. Should discuss this.

```
df <- filter(imdb_details_extd3, type == "Movie") %>%
  na.omit()

data_split <- initial_split(df)
data_train <- training(data_split)
data_test <- testing(data_split)

# build recipe (just instructions)
model_rec <- recipe(gpm ~ budget_adj+runtime+genres+rating+dir_pop_fac+release_period+
  runtime+co_size+star_power+wr_pop,
  data = data_train) %>%
```

```

step_other(genres, threshold = 0.03) %>%
step_unknown(genres) %>%
step_other(rating, threshold = 0.05) %>%
step_unknown(rating) %>%
step_dummy(all_nominal(), -all_outcomes())

# prep actually uses the data
model_prep <- prep(model_rec)
juiced <- juice(model_prep)

# run the below to check the step_other results (doesn't work if step_dummy already used)
# juiced %>% count(genres, sort = T)

# report details
summary(model_rec)

## # A tibble: 10 x 4
##   variable      type    role    source
##   <chr>         <chr>  <chr>   <chr>
## 1 budget_adj    numeric predictor original
## 2 runtime       numeric predictor original
## 3 genres        nominal predictor original
## 4 rating        nominal predictor original
## 5 dir_pop_fac    numeric predictor original
## 6 release_period nominal predictor original
## 7 co_size       numeric predictor original
## 8 star_power     numeric predictor original
## 9 wr_pop        numeric predictor original
## 10 gpm          numeric outcome  original

# build model for tuning
rf_bo_profit_spec <- rand_forest(
  mtry = tune(),
  trees = tune()
) %>%
  set_mode("regression") %>%
  set_engine("ranger")

rf_tune_wf <- workflow() %>%
  add_recipe(model_rec) %>%
  add_model(rf_bo_profit_spec)

# create a set of cross-validation resamples to use for tuning
data_folds <- vfold_cv(data_train)


# choose 10 grid points automatically
rf_tune_res <- tune_grid(
  rf_tune_wf,
  resamples = data_folds,
  grid = 10
)

## i Creating pre-processing data to finalize unknown parameter: mtry

```

```
### rmse plot for tuning mtry and number of trees
```

```
rf_tune_res %>%
  collect_metrics() %>%
  filter(.metric == "rmse") %>%
  dplyr::select(mean, trees, mtry) %>%
  pivot_longer(trees:mtry,
               values_to = "value",
               names_to = "parameter")
) %>%
  ggplot(aes(value, mean, color = parameter)) +
  geom_point(show.legend = FALSE) +
  facet_wrap(~parameter, scales = "free_x") +
  labs(x = NULL, y = "rmse")
```




```
### rsq plot for tuning mtry and number of trees
```

```
rf_tune_res %>%
  collect_metrics() %>%
  filter(.metric == "rsq") %>%
  dplyr::select(mean, trees, mtry) %>%
```

```

pivot_longer(trees:mtry,
              values_to = "value",
              names_to = "parameter"
) %>%
ggplot(aes(value, mean, color = parameter)) +
geom_point(show.legend = FALSE) +
facet_wrap(~parameter, scales = "free_x") +
labs(x = NULL, y = "Rsqr")

```



```
## looks like 3 for mtry and 1000-1500 for trees could work best
```

```
### taking a closer look now
```

```

rf_grid <- grid_regular(
  mtry(range = c(2, 8)),
  trees(range = c(800, 1200)),
  levels = 5
)

```




```

rf_res <- tune_grid(
  rf_tune_wf,
  resamples = data_folds,
  grid = rf_grid
)

## rmse plot for tuning mtry and number of trees
rf_res %>%
  collect_metrics() %>%
  filter(.metric == "rmse") %>%
  mutate(trees = factor(trees)) %>%
  ggplot(aes(mtry, mean, color = trees)) +
  geom_line(alpha = 0.5, size = 1.5) +
  geom_point() +
  labs(y = "rmse")

```



```


## rsq plot for tuning mtry and number of trees
rf_res %>%
  collect_metrics() %>%

```

```

filter(.metric == "rsq") %>%
mutate(trees = factor(trees)) %>%
ggplot(aes(mtry, mean, color = trees)) +
geom_line(alpha = 0.5, size = 1.5) +
geom_point() +
labs(y = "rsq")

```



```
## looks like 4 for mtry and 1000 for trees is optimal
```

```
## build rf model with tuned params
```

```

final_rf <- rand_forest(
  mtry = 4,
  trees = 1000,
) %>%
  set_mode("regression") %>%
  set_engine("ranger")

```


```
# checking out importance plots
```

```
final_rf %>%
```

```

set_engine("ranger", importance = "permutation") %>%
fit(gpm ~ .,
    data = juice(model_prep)
) %>%
vip(geom = "point")

```



```
### view metrics
```

```

final_wf <- workflow() %>%
  add_recipe(model_rec) %>%
  add_model(final_rf)

```

```

final_res <- final_wf %>%
  last_fit(data_split)

```

```

final_res %>%
  collect_metrics()

```

```
## # A tibble: 2 x 4
```

```
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard     0.290 Preprocessor1_Model1
## 2 rsq     standard     0.00469 Preprocessor1_Model1
```

Compare Box Office Profit Random Forest with Linear Model

```
# build recipe (just instructions)
lm_model_rec <- recipe(gpm ~ budget_adj+runtime+dir_pop_fac+release_period+
  runtime+co_size+star_power+wr_pop,
  data = data_train)

## build lm model with tuned params
final_lm <- linear_reg() %>%
  set_engine("lm") %>%
  set_mode("regression")

lm_final_wf <- workflow() %>%
  add_recipe(lm_model_rec) %>%
  add_model(final_lm)

lm_final_res <- lm_final_wf %>%
  last_fit(data_split)

lm_final_res %>%
  collect_metrics()

## # A tibble: 2 x 4
##   .metric .estimator .estimate .config
##   <chr>   <chr>       <dbl> <chr>
## 1 rmse    standard     0.290 Preprocessor1_Model1
## 2 rsq     standard     0.00230 Preprocessor1_Model1
```